


Process Management in Linux

Dr.Hashamdar

Massachusetts Institute of Technology
(MIT)

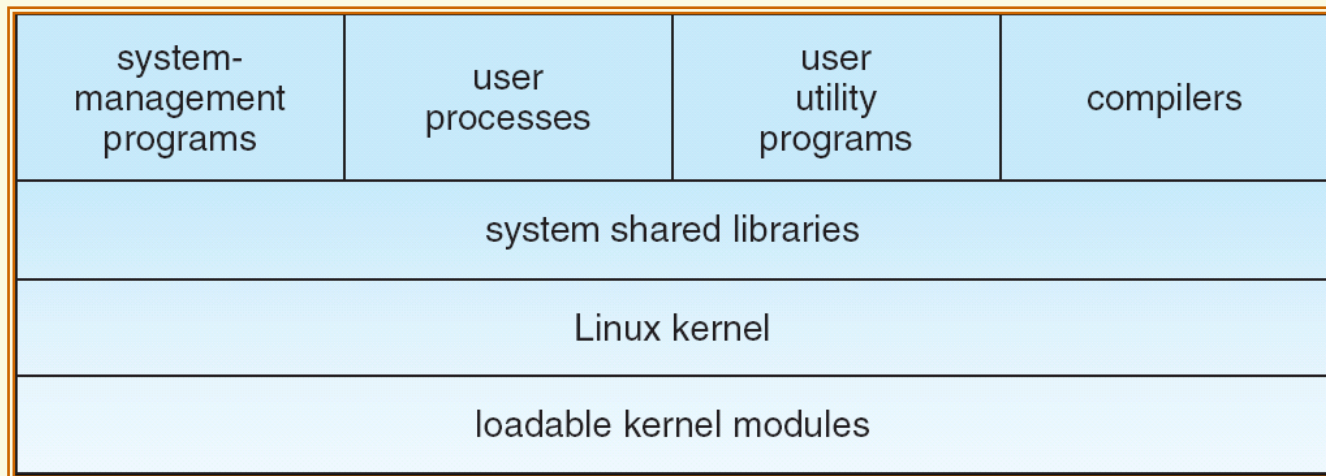


The Linux Operating System




The Linux kernel tracks:

- The location of the process execution (context)
- Which files the process is accessing or has access to
- What users and groups the process belongs to (credentials)
- current directory for the process is
- memory space the process has access to and how it uses it.



What is a process?



What is a process?

Processes are not just **programs**

- Program instructions plus other components as needed (primarily data)

Sometimes called **jobs** or **tasks**

Process Management

- ❏ UNIX process management separates the creation of processes and the running of a new program into two distinct operations.
 - The fork system call creates a new process
 - A new program is run after a call to exec
- ❏ Under UNIX, a process encompasses all the information that the operating system must maintain to track the context of a single execution of a single program

Process Management



Under Linux, process properties fall into 3 groups

- (i) Process's Identity
- (ii) Process's Environment
- (iii) Process's Context

(i) Process Identity



Three different attributes for a *process identity*

- a) Process ID (PID)
- b) Credentials
- c) Personality

a) Process ID (PID)

- 📄 PID is a unique integer
- 📄 PID cannot be changed by the process
- 📄 Kernel uses PID to track a process
 - Tracks exit status
 - Relationship to other processes

a) Process ID (PID)

 When process exits

- PID is kept until it can be safely discarded.
- PID that is kept in an active status is called a zombie.
- All child processes (now called orphan processes) become child processed of the init process.

(i) Process Identity



Three different attributes for a *process identity*

a) Process ID (PID)

b) Credentials

c) Personality

b) Credentials

☞ Are used to insure security

☞ Made up of

- users
- and groups

☞ User ID (UID) and Group ID (GID)

- either one can be set within a limit
- symbolic user and group names that are mapped to a unique integer value (usually positive).

(i) Process Identity



Three different attributes for a *process identity*

a) Process ID (PID)

b) Credentials

c) Personality

c) Personality

- Personality identifiers allow slight modifications to the semantics of certain system calls.
- “Personalities are primarily used by emulation libraries to request that system calls can be compatible with certain specific flavors of unit.” (Galvin, 710)

(i) Process Identity



Three different attributes for a *process identity*

- a) Process ID (PID)
- b) Credentials
- c) Personality

Process Management



Under Linux, process properties fall into 3 groups

- (i) **Process's Identity**
- (ii) Process's Environment
- (iii) Process's Context



(ii) Process Environment



The process's environment is inherited from its parent, and is composed of **2 null-terminated vectors**:

- The argument vector lists the command-line arguments used to invoke the running program; conventionally starts with the name of the program itself
- The environment vector is a list of “**NAME=VALUE**” pairs that associates named environment variables with arbitrary textual values

(ii) Process Environment

-  Passing environment variables among processes and inheriting variables by a process's children are flexible means of passing information to components of the user-mode system software
-  The environment-variable mechanism provides a customization of the operating system that can be set on a per-process basis, rather than being configured for the system as a whole

Process Management






Under Linux, process properties fall into 3 groups

- (i) Process's Identity
- (ii) Process's Environment
- (iii) Process's Context

(iii) Process Context

- ❏ The (constantly changing) state of a running program at any point in time
- ❏ The **scheduling context** is the most important part of the process context; it is the information that the scheduler needs to suspend and restart the process
- ❏ The kernel maintains **accounting** information about the resources currently being consumed by each process, and the total resources consumed by the process in its lifetime so far

(iii) Process Context

-  The **file table** is an array of pointers to kernel file structures
 - When making file I/O system calls, processes refer to files by their index into this table
-  The **signal-handler table** defines the routine in the process's address space to be called when specific signals arrive
-  The **virtual-memory context** of a process describes the full contents of the its private address space

Process Management



Under Linux, process properties fall into 3 groups

- (i) Process's Identity
- (ii) Process's Environment
- (iii) Process's Context

Processes and Threads

- Linux uses the same internal representation for processes and threads;
- a thread is simply a new process that happens to share the same address space as its parent


Processes and Threads

- 📄 A distinction is only made when a new thread is created by the **clone** system call
 - **Fork** creates a new process with its own entirely new process context
 - **clone** creates a new process with its own identity, but that is allowed to share the data structures of its parent
- 📄 Using **clone** gives an application fine-grained control over exactly what is shared between two threads

Organization of Table of Processes

- ❏ Each process is referenced by **descriptor**
 - Describes process attributes together with information needed to manage process
- ❏ Kernel dynamically allocates these descriptors when processes begin execution
- ❏ All process descriptors are organized in **doubly linked list**
- ❏ Scheduler used **Macro instructions** to manage and update process descriptor lists as needed

Process Synchronization

 To allow two processes to synchronize with each other, Linux provides:

- **Wait queue:** Linked circular list of process descriptors
- **Semaphores:** Used to solve problems of mutual exclusion and problems of producers and consumers
 - In Linux they contain 3 fields:
 - Semaphore counter
 - Number of waiting processes
 - List of processes waiting for semaphore